

# STATS 551 - HW3

Zhen Qin

1.

The model is  $y_j|\theta_j \sim \text{binomial}(n_j, \theta_j)$ .  $n_j$  is the total number of trials.  $\theta_j \sim \text{beta}(\alpha, \beta)$ ,  $j = 1, \dots, 68$ . Suppose that  $(\alpha, \beta)$  is given. The joint posterior distribution is  $p(\theta_1, \dots, \theta_n | \alpha, \beta, y_1, \dots, y_n) \propto \prod_j \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta_j^{\alpha-1} (1 - \theta_j)^{\beta-1} \theta_j^{y_j} (1 - \theta_j)^{n_j - y_j} = \prod_j \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta_j^{\alpha+y_j-1} (1 - \theta_j)^{\beta+n_j-y_j-1}$ .

Given  $(\alpha, \beta) = (1.4, 8.6)$ ,  $\theta_j$  are all independent, so visualization of one parameter is enough. In fact, after integration we know  $\theta_i | \alpha, \beta, \text{obs} \sim \text{beta}(\alpha + y_i, \beta + n_i - y_i)$ .

```
library(ggplot2)
library(gridExtra)
library(tidyr)


# load data
y <- c(0,0,0,0,0,0,0,0,0,0,0,0,1,1,1,1,1,1,2,2,2,2,2,2,2,2,
2,1,5,2,5,3,2,7,7,3,3,2,9,10,4,4,4,4,4,4,4,10,4,4,4,5,
11,12, 5,5,6,5,6,6,6,6,16,15,15,9,4)
n <- c(20,20,20,20,20,20,19,19,19,19,18,17,20,20,20,20,19,19,18,18,25,24,
23,20,20,20,20,20,20,10,49,19,46,27,17,49,47,20,20,13,48,50,20,
20,20,20,20,20,20,48,19,19,19,22,46,49,20,20,23,19,22,20,20,20,
52,46,47,24,14)

# create the grid
x <- seq(0.0001, 0.9999, length.out = 1000)

# beta distribution
bdens <- function(n, y, x)
  dbeta(x, y+1.4, n-y+8.6)
df_sep <- mapply(bdens, n, y, MoreArgs = list(x = x)) %>%
  as.data.frame() %>% cbind(x) %>% gather(ind, p, -x)

# plot the separate model
plot_sep <- ggplot(data = df_sep) +
  geom_line(aes(x = x, y = p, group = ind)) +
  labs(x = expression(theta), y = '', title = 'Joint Posterior Distribution', color = '') +
  scale_y_continuous(breaks = NULL) +
  theme(legend.background = element_blank(), legend.position = c(0.8,0.9))
# The last one is for emphasize colored red
plot_sep
```

## Joint Posterior Distribution



2.

According to the distributions,  $E(y_i|\alpha, \beta) = E(E(y_i|\alpha, \beta, \theta_i)) = E(n_i\theta_i) = n_i\frac{\alpha}{\alpha+\beta}$ ,  $E(y_i^2|\alpha, \beta) = E(E(y_i^2|\alpha, \beta, \theta_i)) = E(n_i\theta_i(1 - \theta_i) + (n_i\theta_i)^2) = n_iE(\theta_i) + (n_i^2 - n_i)E(\theta^2) = n_i\frac{\alpha}{\alpha+\beta} + (n_i^2 - n_i)(\frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)} + (\frac{\alpha}{\alpha+\beta})^2)$ . Let the moment of population equal the moment of samples, then  $1/68 \times \sum_i y_i/n_i = \frac{\hat{\alpha}}{\hat{\alpha}+\hat{\beta}}$ ,  $(1/68 \times \sum_i (y_i^2 - y_i)/(n_i^2 - n_i)) = \frac{\hat{\alpha}\hat{\beta}}{(\hat{\alpha}+\hat{\beta})^2(\hat{\alpha}+\hat{\beta}+1)} + (\frac{\hat{\alpha}}{\hat{\alpha}+\hat{\beta}})^2$ , I solved these two equations and get  $\hat{\alpha} = 3.17$ ,  $\hat{\beta} = 18.92$ , so the posterior distribution is  $\theta_i \sim \text{beta}(3.17 + y_i, 18.92 + n_i - y_i)$ .

In question 1,  $\alpha, \beta$  do not contain information from data, so posterior densities differ a lot. In question 2,  $\alpha, \beta$  are the estimators by MOM, so they are close to the true parameters and posterior densities are more stable.


```
tmp=mean((y^2-y)/(n^2-n))/(mean(y/n))^2-1
tmp=tmp*mean(y/n)/(1-mean(y/n))
tmp=1/tmp-1
alphahat=tmp*mean(y/n)
betahat=tmp-alphahat

# beta distribution
bdens <- function(n, y, x)
  dbeta(x, y+alphahat, n-y+betahat)
df_sep <- mapply(bdens, n, y, MoreArgs = list(x = x)) %>%
  as.data.frame() %>% cbind(x) %>% gather(ind, p, -x)

# plot the separate model
```

```
plot_sep <- ggplot(data = df_sep) +
  geom_line(aes(x = x, y = p, group = ind)) +
  labs(x = expression(theta), y = '', title = 'Joint Posterior Distribution', color = '') +
  scale_y_continuous(breaks = NULL) +
  theme(legend.background = element_blank(), legend.position = c(0.8,0.9))
# The last one is for emphasize colored red
plot_sep
```

Joint Posterior Distribution



3.

In this question, I use Jacobian determinant to calculate the density of transformations of Random Variables. Suppose  $x = \frac{\alpha}{\alpha + \beta}$ ,  $y = (\alpha + \beta)^{-1/2}$ , then  $\alpha = x/y^2$ ,  $\beta = (1 - x)/y^2$  and  $(x, y)$  obeys uniform distribution. Then  $p(\alpha, \beta) \propto 1 \times |J|^{-1} \propto 1/|\frac{2x}{y^5} + \frac{2-2x}{y^5}| \propto y^5 = (\alpha + \beta)^{-5/2}$ , where  $|J|$  is the Jacobian determinant.

As for the log transformed scale,  $\log(\alpha/\beta) = \log(x/(1-x))$ ,  $\log(\alpha + \beta) = -2\log y$ , so  $p(\log(\alpha/\beta), \log(\alpha + \beta)) \propto 1 \times |J|^{-1} \propto 1/|\frac{1}{x(1-x)} \times \frac{-2}{y}| \propto x(1-x)y \propto \alpha\beta(\alpha + \beta)^{-5/2}$ , where  $|J|$  is the Jacobian determinant.

4.

The joint posterior distribution is  $p(\alpha, \beta, \theta_1, \dots, \theta_n | y_1, \dots, y_n) \propto p(\alpha, \beta) \prod_j \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta_j^{\alpha-1} (1 - \theta_j)^{\beta-1} \theta_j^{y_j} (1 - \theta_j)^{n_j - y_j} = p(\alpha, \beta) \prod_j \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta_j^{\alpha + y_j - 1} (1 - \theta_j)^{\beta + n_j - y_j - 1}$ . Suppose that  $(\alpha, \beta)$  obey a noninformative hyperprior distribution, i.e.  $p(\alpha, \beta) \propto (\alpha + \beta)^{-5/2}$ . The joint posterior distribution is  $p(\alpha, \beta, \theta_1, \dots, \theta_n | y_1, \dots, y_n) \propto (\alpha + \beta)^{-5/2} \prod_j \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta_j^{\alpha-1} (1 - \theta_j)^{\beta-1} \theta_j^{y_j} (1 - \theta_j)^{n_j - y_j} = (\alpha + \beta)^{-5/2} \prod_j \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta_j^{\alpha + y_j - 1} (1 - \theta_j)^{\beta + n_j - y_j - 1}$ .

On the original scale the marginal distribution is  $p(\alpha, \beta | obs) = \int p(\alpha, \beta, \theta_1, \dots, \theta_{68} | y_1, \dots, y_{68}) d\theta_1 \dots d\theta_{68} \propto (\alpha + \beta)^{-5/2} \prod_j \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \frac{\Gamma(\alpha + y_j)\Gamma(\beta + n_j - y_j)}{\Gamma(\alpha + \beta + n_j)}$ .

On the log transformed scale the marginal distribution is  $p(\log \frac{\alpha}{\beta}, \log(\alpha + \beta) | obs) = \int p(\log \frac{\alpha}{\beta}, \log(\alpha + \beta), \theta_1, \dots, \theta_{68} | y_1, \dots, y_{68}) d\theta_1 \dots d\theta_{68} \propto \alpha\beta(\alpha + \beta)^{-5/2} \prod_j \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \frac{\Gamma(\alpha + y_j)\Gamma(\beta + n_j - y_j)}{\Gamma(\alpha + \beta + n_j)}$ .

## 5.

Work with the transformed scale.


```
library(mvtnorm)
mcmc_MH <- function(Niter, proposal_sigma, mu_init, LL){
  mu_new <- mu_init
  d <- length(mu_new)
  samples_MH <- mu_new
  accept_MH <- 0
  for(iter in 1:Niter){
    if(d > 1){
      propose_MH_mu <- as.numeric(rmvnorm(1, mean = mu_new,
                                          sigma = diag(rep(proposal_sigma^2, d))))
    }
    if(d == 1){
      propose_MH_mu <- rnorm(1, mean = mu_new, sd = proposal_sigma)
    }
    rtemp <- LL(propose_MH_mu) - LL(mu_new)
    if(runif(1) < exp(rtemp)){
      mu_new <- propose_MH_mu
      accept_MH <- accept_MH + 1
    }
    samples_MH <- rbind(samples_MH, mu_new)
  }
  return(list(samples = samples_MH, acceptance = accept_MH /Niter))
}

LL <- function(x){
  a = exp(sum(x))/(exp(x[1])+1)
  b = exp(x[2])/(exp(x[1])+1)
  ll1=log(a)+log(b)-2.5*log(a+b)
  ll2=log(gamma(a+y))+log(gamma(b+n-y))-log(gamma(a+b+n))+
    log(gamma(a+b))-log(gamma(a))-log(gamma(b))
  ll1+sum(ll2)
}

trans_inverse = function(x){
  a = exp(sum(x))/(exp(x[1])+1)
  b = exp(x[2])/(exp(x[1])+1)
  c(a,b)
}


Niter = 1000
proposal_sigma = 0.1 # try other values
mu_init = c(-1,1) # try other values
samplesMH = mcmc_MH (Niter, proposal_sigma, mu_init, LL)
```

```
absamples=data.frame(alpha=exp(rowSums(samplesMH$samples))/(exp(samplesMH$samples[,1])+1),
                     beta=exp(samplesMH$samples[,2])/(exp(samplesMH$samples[,1])+1))
plot(apply(absamples, 2, jitter), type = 'l', xlab = '', ylab = '')
points(t(trans_inverse(mu_init)), cex = 2, col = 'red')
```



6.

```
# beta distribution
samplesttheta=matrix(rep(0,68*nrow(absamples)),nrow = nrow(absamples))
for(i in 1:nrow(absamples)){
  for(j in 1:68){
    samplesttheta[i,j]=rbeta(1,absamples[i,1]+y[j],absamples[i,2]+n[j]-y[j])
  }
}
boxplot(samplesttheta)
```



7.

```
library(mvtnorm)
mcmc_MHG <- function(Niter, proposal_sigma, mu_init, LL){
  mu_new <- mu_init
  d <- length(mu_new)
  samples_MH <- mu_new
  thetas=LL(samples_MH)$thetas
  accept_MH <- 0
  for(iter in 1: Niter){
    if(d > 1){
      propose_MH_mu <- as.numeric(rmvnorm(1, mean = mu_new,
                                          sigma = diag(rep(proposal_sigma^2, d))))
    }
    if(d == 1){
      propose_MH_mu <- rnorm(1, mean = mu_new, sd = proposal_sigma)
    }
    rtemp <- LL(propose_MH_mu)$ll - LL(mu_new)$ll
    if(runif(1) < exp(rtemp)){
      mu_new <- propose_MH_mu
      accept_MH <- accept_MH + 1
    }
    samples_MH <- rbind(samples_MH, mu_new)
    thetas=rbind(thetas, LL(propose_MH_mu)$thetas)
  }
}
```


```

    }
    absamples=data.frame(alpha=exp(rowSums(samplesMH$samples))/(exp(samplesMH$samples[,1])+1),
                        beta=exp(samplesMH$samples[,2])/(exp(samplesMH$samples[,1])+1))
    return(list(samples = cbind(absamples,thetas), acceptance = accept_MH /Niter))
}

LL = function(x){
  a = exp(sum(x))/(exp(x[1])+1)
  b = exp(x[2])/(exp(x[1])+1)
  samplestheta=1:68
  for(i in 1:68){
    samplestheta[i]=rbeta(1, a+y[i],b+n[i]-y[i])
  }
  ll1=log(a)+log(b)-2.5*log(a+b)
  ll2=log(gamma(a+y))+log(gamma(b+n-y))-log(gamma(a+b+n))+
    (a+y-1)*log(samplestheta)+(b+n-y-1)*log(1-samplestheta)
  ll=ll1+sum(ll2)
  return(list(ll=ll,thetas=samplestheta))
}


Niter = 1000
proposal_sigma = 0.1 # try other values
mu_init = c(-1,1) # try other values
samplesMH = mcmc_MHG (Niter, proposal_sigma, mu_init, LL)
plot(apply(samplesMH$samples[,1:2], 2, jitter), type = 'l', xlab = '', ylab = '')
points(t(trans_inverse(mu_init)), cex = 2, col = 'red')

```



```
boxplot(samplesMH$samples[,3:ncol(samplesMH$samples)])
```






8.

Work with the original scale.


```
library(rstan)
library(rstudioapi)
rat_data=list(J = 68, y=y,n=n)
fit <- stan(file = 'rat.stan', data = rat_data,
            iter = 1000, chains = 1)
```

```
##
## SAMPLING FOR MODEL 'rat' NOW (CHAIN 1).
##
## Gradient evaluation took 0 seconds
## 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Adjust your expectations accordingly!
##
##
## Iteration: 1 / 1000 [ 0%] (Warmup)
## Iteration: 100 / 1000 [ 10%] (Warmup)
## Iteration: 200 / 1000 [ 20%] (Warmup)
## Iteration: 300 / 1000 [ 30%] (Warmup)
## Iteration: 400 / 1000 [ 40%] (Warmup)
## Iteration: 500 / 1000 [ 50%] (Warmup)
## Iteration: 501 / 1000 [ 50%] (Sampling)
## Iteration: 600 / 1000 [ 60%] (Sampling)
```

```
## Iteration: 700 / 1000 [ 70%] (Sampling)
## Iteration: 800 / 1000 [ 80%] (Sampling)
## Iteration: 900 / 1000 [ 90%] (Sampling)
## Iteration: 1000 / 1000 [100%] (Sampling)
##
## Elapsed Time: 0.384 seconds (Warm-up)
##               0.393 seconds (Sampling)
##               0.777 seconds (Total)
la <- extract(fit, permuted = TRUE)
plot(apply(cbind(la$alpha, la$beta), 2, jitter), type = 'l', xlab = '', ylab = '')
```



```
boxplot(la$theta)
```



9.

```
mean(samplesMH$samples[,1])
```

```
## [1] 2.833363
```

```
mean(samplesMH$samples[,1]<0.2*(samplesMH$samples[,1]+samplesMH$samples[,2]))
```

```
## [1] 0.991009
```

10.

```
mean(la$alpha)
```

```
## [1] 4.119804
```

```
mean(la$alpha<0.2*(la$alpha+la$beta))
```

```
## [1] 1
```

11.

The flat priors is not the proper posterior. because if this prior is used, then  $\alpha + \beta \rightarrow \infty$ . In this case the marginal distribution is  $p(\alpha, \beta | obs) \propto \prod_j \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} \frac{\Gamma(\alpha+y_j)\Gamma(\beta+n_j-y_j)}{\Gamma(\alpha+\beta+n_j)}$ . From the plot below we know when  $\alpha, \beta$

increases by  $p$  times ( $p > 1$ ), the log-likelihood increases, thus the integration of the density does not exist. So it is not a proper prior.

```
ll = function(x){  
  a=x[,1]  
  b=x[,2]  
  y=4  
  n=14  
  log(gamma(a+y))+log(gamma(b+n-y))-log(gamma(a+b+n))+  
  log(gamma(a+b))-log(gamma(a))-log(gamma(b))  
}  
x=data.frame(a=(1:68)/3,b=1:68)  
plot(x$a+x$b,ll(x),xlab = 'alpha+beta',ylab = 'log-density')
```

